Rishi Guptha Mankala

United States

934-263-3087 | rishiguptha.mankala@gmail.com | Linkedin

Github | Portfolio

WORK EXPERIENCE

Nevara AI | *AI/ML Engineering Intern* — **Jun 2025 - Aug 2025**

- Developed an AI sales-call auditing MVP using Next.js/TypeScript, Supabase, OpenAI, and Deepgram for a coach overseeing 100+ reps. Automated notes-to-report in 90–140 seconds, reducing per-rep review time by approximately 75%, and achieved a satisfaction rating of 3.5–3.75 out of 5 in a pilot with 3 reps across 3 calls
- Owned CI/CD and serverless delivery with GitHub Actions to Vercel and FastAPI on AWS Lambda (Mangum, Docker). Tuned memory and reserved/provisioned concurrency to stabilize p95 latency at 30s (transcribe), 30–50s (notes), and 60–90s (report).
- Shipped timestamp-synced media player, transcript, and notes (HLS.js, SSE, React Query). Added click-to-seek and auto-highlighting that improved audit navigation throughput by 2x and reduced context switching.
- Created Next.js API routes for transcription, ai-report, gong-scrape, fathom with SSE progress, XML validation, and versioned Supabase storage using Auth/Postgres/Storage via Prisma. Integrated Gong and Fathom, leading AI-report prompt and context engineering to maintain costs under $0.5 per call and reduce retries by approximately 30%, improving system efficiency and cost-effectiveness

Stony Brook University | *Data Science Researcher* — **Jan 2025 - May 2025**

- Built Python-SQL pipelines for S&P 500 time-series data and trained LSTM, CNN, and SVM models using TensorFlow/PyTorch, improving model accuracy by 20% and risk scoring by 15% through Bayesian inference and backtesting.
- Delivered dashboards (Tableau, Python) integrating predictive outputs and risk metrics, enabling 30% faster portfolio decisions and surfacing volatility trends via SVM optimization and feature extraction.

Kimberly Clark | *Data Engineering Intern* — **Jan 2023 - Jul 2023**

- Automated metadata tagging and validation in Collibra using Python/PostgreSQL workflows, earning 7 certifications and boosting governance team productivity by 25% while reducing manual effort by 30%.
- Migrated 10K+ assets into Collibra via Excel-based ETL with 100% data integrity, resolved 15% pipeline inconsistencies, and built Power BI/Tableau dashboards to visualize KPIs and reduce reporting errors by 18%.
- Designed and deployed 5+ advanced Collibra workflows to support real-time data quality checks and lineage validation, enabling faster stakeholder reporting and improved oversight across governance domains.

Matchday AI | *Data Scientist Intern* — **Jan 2022 - Dec 2022**

- Built computer vision pipelines with OpenCV and deep learning to detect field lines, ball position, and player movement from broadcast footage, achieving 90%+ event tracking accuracy.
- Applied Homography Transformation to convert dynamic camera feeds into 2D field projections, powering spatial analytics for tactical coaching insights.
- Developed a full-stack annotation platform (Flask, JS, SQLite), reducing manual labeling time by 40% and enabling deployment of data-driven dashboards for 6+ ISL teams via Star Sports.

PROJECTS

NutriSmart: RAG-Driven Insights for Fitness and Nutrition | *Academic*

- Delivered 30% faster fitness insights and 40% lower latency by deploying a RAG service (Sentence Transformers + FAISS) containerized with Docker on AWS EC2/FastAPI, indexing 1.2K+ papers.

NYC Taxi Insights | *Academic*

- Improved fare/duration accuracy 30% and query performance 35% by engineering a GCP ETL for 10M+ rides, modeling star-schema marts, and training LightGBM with BERT-derived features for dynamic pricing.

Sales Insights Dashboard | *Individual*

- Raised sales-forecast accuracy 30% and cut prediction turnaround 15% by building a Power BI hub over CRM/ERP using SQL + DAX, enabling faster pipeline planning and territory targeting.

Spotify Playlist Analytics | *Individual*

- Enabled daily KPIs by designing a serverless ETL (AWS Lambda/Glue/S3) for 18K+ playlist logs and optimizing Athena SQL to reduce processing time 40%, informing genre strategy and retention.

EDUCATION

Stony Brook University | *Masters of Science, Data Science* — **Stony Brook, NY**

SKILLS

- **Programming Languages:** Python, Java, SQL, TypeScript/JavaScript, R
- **Data Science & Machine Learning:** PyTorch, TensorFlow, scikit-learn, Pandas, NumPy, NLP, ASR (Deepgram), Large Language Models (OpenAI API), Prompt & Context Engineering, Statistical Analysis/Modeling, Deep Learning
- **Cloud Platforms:** AWS (Lambda, S3, EC2, Glue, Redshift), GCP (BigQuery, Cloud Storage, Vertex AI), Vercel, Supabase (Auth/Postgres/Storage)
- **Data Platforms & Pipelines:** Snowflake, Apache Spark, Airflow (Astronomer), dbt, Data Modeling, PostgreSQL
- **Backend & Web:** Next.js (React), FastAPI, RESTful APIs
- **Data Visualization:** Tableau, Power BI, Looker Studio, AWS QuickSight, Matplotlib
- **DevOps & Tools:** Docker, Linux, Git/GitHub, GitHub Actions (CI/CD), Microsoft Excel, Google Sheets, Collibra